


Operating Systems

Race Conditions

Hongliang Liang, BUPT

Sep. 2023

Communication Across Threads



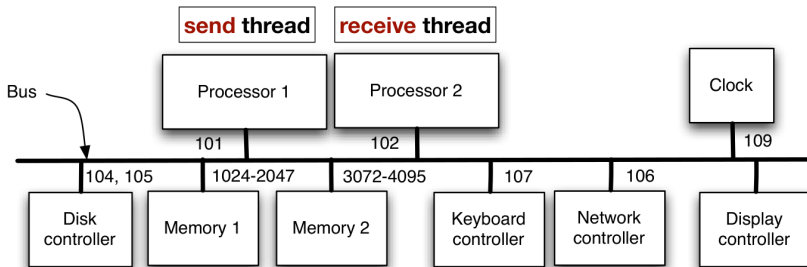
Communication Across Threads: Primitives

*The OS may provide the following interfaces for **send** and **receive** with bounded buffers:*

- `send(message)`: if there is room in the bounded buffer, insert message in the buffer. If not, **stop the calling thread and wait** until there is room.
- `receive()`: if there is a message in the bounded buffer, return the message to the calling thread. If not, **stop the calling thread and wait** until another thread sends a message to buffer.

First assumption: one CPU per thread

For now, let's assume that there is an available physical CPU for each thread, so we don't need to worry about multiplexing threads on the same CPU



The Producer-Consumer Problem

- The problem of sharing a bounded buffer between two threads is an instance of the producer-consumer problem
 - The producer needs to first add a message to the shared buffer before the consumer can remove it
 - The producer needs to wait for the consumer to catch up when the buffer fills up
- Let's try to implement `send()` and `receive()`
 - We can alternatively call them `producer()` and `consumer()`

First implementation of send() and receive()

```
message buffer[N]
int in = 0, out = 0
send(message msg)
    while in - out == N do nothing
    buffer[in % N] = msg
    in = in + 1
    return
message receive()
    while in == out do nothing
    msg = buffer[out % N]
    out = out + 1
    return msg
```

Assumptions revisited

- There is only one sending thread and one receiving thread
 - Only one thread updates each shared variable
 - Only the receiving thread updates **out**
 - Only the sending thread updates **in**
- **One writer principle:** If each variable has only one writer, coordination becomes easier.

What if we allow several senders?

- Each sender has its own CPU, but progresses at different paces

A buffer is empty fill entry 0 in=1

----->

B buffer is empty fill entry 0 in = 2

----->

- This is called a **race condition**, since it depends on the exact timing of two threads

Another race condition

```
in = in + 1
```

```
1 LOAD in, R0
```

```
2 ADD R0, 1
```

```
3 STORE R0, in
```

```
read 0
```

A 1 2 3 in = 1

----->

```
read 0
```

B 1 2 3 in = 2

----->

Very difficult to debug, as it may rarely occur, and hard to reproduce

How do we fix race conditions?

- Race conditions occur whenever the output depends on the precise execution order of threads
- How do we systematically avoid race conditions?

How do we fix race conditions?

- Intuitively, we need to make threads coordinate with each other to ensure **mutual exclusion** in accessing **critical sections** of code
 - In this case, a critical section defines a multi-step operation on shared data (e.g., variables) that needs to become “before-or-after” **atomic actions**
- We need a **lock**!

Shared locks to achieve mutual exclusion

- A **lock** is a shared variable that acts as a flag to coordinate usage of other shared variables
- We introduce two new primitives to work with locks
 - **acquire()** and **release()**
 - Now a thread can acquire a lock, hold it for a while, and then release it
 - When a thread is holding a lock, other threads that attempt to acquire that same lock will wait until the first thread releases the lock
 - Analogy: an exclusive restaurant has a table and a chair. when a customer is served, others must wait outside.
- By surrounding multi-step operations involving shared variables with **acquire()** and **release()**, we make sure a multistep operation behaves like a single-step operation
alternatively called **lock()** and **unlock()** (in BLITZ)

Second implementation of send() and receive()

```
message buffer[N]
int in = 0, out = 0
lock buffer_lock = UNLOCKED // add
send(message msg)
    acquire(buffer_lock) // add
    while in - out == N do nothing
    buffer[in % N] = msg
    in = in + 1
    release(buffer_lock) // add
    return
message receive()
    acquire(buffer_lock) // add
    while in == out do nothing
    msg = buffer[out % N]
    out = out + 1
    release(buffer_lock) // add
    return msg
```

Second implementation of send() and receive()

```
message buffer[N]
int in = 0, out = 0
lock buffer_lock = UNLOCKED // add
send(message msg)
    acquire(buffer_lock) // add
    while in - out == N do nothing
    buffer[in % N] = msg
    in = in + 1
    release(buffer_lock) // add
    return
message receive()
    acquire(buffer_lock) // add
    while in == out do nothing
    msg = buffer[out % N]
    out = out + 1
    release(buffer_lock) // add
    return msg
```

Correct? e.g., sender holding the lock faces the full buffer.

Third implementation of send() and receive()

```
message buffer[N]
int in = 0, out = 0
lock buffer_lock = UNLOCKED
send(message msg)
    acquire(buffer_lock)
    while in - out == N do
        release(buffer_lock) //
        acquire(buffer_lock) //
    buffer[in % N] = msg
    in = in + 1
    release(buffer_lock)
    return
message receive()
    acquire(buffer_lock)
    while in == out do
        release(buffer_lock) //
        acquire(buffer_lock) //
    msg = buffer[out % N]
    out = out + 1
    release(buffer_lock)
    return msg
```

Implementing acquire() and release()

- A correct implementation must enforce the “single-acquire” protocol

Several threads may attempt to acquire the lock at the same time, but only one should succeed

- Consider the following implementation —

Implementing acquire() and release()

```
struct lock  
    int state
```

```
acquire(lock L)  
    while L.state == LOCKED do nothing  
    L.state = LOCKED
```

```
release(lock L)  
    L.state = UNLOCKED
```

Race condition

A L.state is UNLOCKED L.state = LOCKED

B L.state is UNLOCKED L.state = LOCKED

Why the race condition?

- The faulty **acquire()** has a multi-step operation on a shared variable (the lock), and we must ensure in some way that **acquire()** itself is a before-or-after atomic action
- Once **acquire()** is an atomic action, we can use it to turn arbitrary multi-step operations on shared variables into before-or-after atomic actions
- Did we just go back to the very beginning?

Why the race condition?

- The faulty **acquire()** has a multi-step operation on a shared variable (the lock), and we must ensure in some way that **acquire()** itself is a before-or-after atomic action
- Once **acquire()** is an atomic action, we can use it to turn arbitrary multi-step operations on shared variables into before-or-after atomic actions
- Did we just go back to the very beginning?
- No, we have actually made progress!
- We reduced a more general problem (making multi-step operations on shared variables before-or-after actions) to a narrower problem (making an operation on a single shared lock a before-or-after action)!

Disabling interrupts?

- Timer interrupts are disabled
- The thread scheduler is not able to run
- No other thread can run on this CPU

Disabling interrupts?

Problems:

- If we have multiple CPUs, threads running on the other processors can enter the critical section even with interrupts disabled
- Is it fine to trust the user threads for disabling interrupts?
- If so, what if they never re-enable the interrupts?

Hardware support: we need it again

Hardware support: the TSL instruction

- Test and Set Lock (TSL) instruction (or called Test and Set) from the hardware, modifies the content of a memory location, returning its old value

TSL(LOCK)	<code>bool tas(bool *lock){</code>
do atomic	<code>bool old = *lock;</code>
1 RX = LOCK	<code>*lock = true;</code>
2 LOCK = LOCKED	<code>return old;</code>
RETURN RX	<code>}</code>

- The bus arbiter in hardware that controls the bus connecting processors to the memory must guarantee —
 - the LOAD (line 1) and STORE (line 2) instructions must execute as before-or-after atomic actions
 - By allowing both to be done in a single clock cycle

Implementing acquire() and release() with TSL

```
acquire(lock L)
    R1 = TSL(L.state)
    while R1 == LOCKED do
        R1 = TSL(L.state)

release(lock L)
    L.state = UNLOCKED
```

Correctness of the solution

- To see that the implementation is correct, we assume L is UNLOCKED
- If a thread calls **acquire(L)**, then after TSL, L is LOCKED and R1 contains UNLOCKED, so the thread acquired the lock
- The next thread that calls acquire(L) sees LOCKED in R1 after TSL, showing that a thread is holding the lock
- The thread that tried to acquire L will **spin** until R1 contains UNLOCKED
- When releasing a lock, no test is needed, an ordinary **STORE** instruction can do the job without creating a race condition

Compare-and-swap lock

```
bool cas(T *ptr, T expt, T new) {  
    if (*ptr == expt) {  
        *ptr = new;  
        return true;  
    }  
    return false;  
}
```

The function compares the value at `*ptr` and if it is equal to `expt` then the value is overwritten with `new`. The function returns `true` if the swap happened.

Compare-and-swap lock

How would you implement the lock acquire operation?

Compare-and-swap lock

How would you implement the lock acquire operation?

```
void acquire_cas(bool *lck) {  
    while (compare_and_swap(lck, false, true)  
           == false);  
}
```

- **Mutual exclusion** from a **critical section** using **shared locks**
 - No concurrent threads are simultaneously in the critical section at the same time
 - Shared locks are usually called **mutex locks** (mutual exclusion locks)
- Two operations to protect the critical section
 - acquire(L): alternatively called lock(L)
 - release(L): alternatively called unlock(L)
- Implementation
 - Using the Test-and-Set-Lock (TSL) or Compare-and-Swap-Lock instruction: an atomic action
- With mutex locks, the producer-consumer problem is solved
 - allowing multiple senders and receivers to access a shared buffer